

Palette: A Decision & Reliability System for AI in Production

The Problem

AI systems fail in production not because models break, but because humans lose alignment. Teams build solutions to problems they misunderstood. Decisions get made without recording why. Work becomes non-restartable when people leave or context is lost. Silent decision debt accumulates not through post-hoc interpretability, but by preserving the reasoning that led to each decision.

It enables safe iteration. Engineers can modify systems confidently because they know which changes are reversible and which require careful review. Product teams can trace outcomes back to decisions. Operators have decision logs to guide incident response.

It scales institutional knowledge. When people transition, the work remains restartable. New team members don't inherit black boxesit structures collaboration so humans and AI can make better decisions together, with full visibility into what was decided and why.

Relevance to Core AI

Core AI's mandate is to turn ambitious research into reliable, productized systems. That requires more than model performancesystems that teams can understand, modify, and trust. Palette is evidence of how I approach that problem: not through abstraction or automation, but through disciplined collaboration and preserved reasoning.

This is the work I do. This is how I operate.